

Class 7

Jimmi Nguyen


```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

```
tmp <- c(rnorm(30,-3), rnorm(30,3))  
x <- cbind(x=tmp, y=rev(tmp))  
head(x)
```

	x	y
[1,]	-2.197415	1.607807
[2,]	-2.037092	1.922347
[3,]	-2.832612	2.585784
[4,]	-2.618156	3.194086
[5,]	-2.852086	4.192869
[6,]	-2.997077	3.633939

Quick plot of x to see the two groups at -3,+3 and +3,-3

```
plot(x)
```



Use `kmeans()` function setting `k` to 2 and `nstart = 20`

```
km = kmeans(x, centers = 2, nstart = 20)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.852485	3.039061
2	3.039061	-2.852485

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 57.89978 57.89978
(between_SS / total_SS = 90.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q. How many points are in each cluster?

km\$size

[1] 30 30

Q. What ‘component’ of your results object details - cluster assignment/membership - cluster center?

km\$cluster


[illegible]

km\$centers

	x	y
1	-2.852485	3.039061
2	3.039061	-2.852485


Q. Plot `x` colored by the `kmeans` cluster assignment and add cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```



Play with kmeans and ask for different number of clusters

```
km = kmeans(x, centers = 4, nstart = 20)
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```



Hierarchical Clustering

This is another very useful and widely employed clustering method which has the advantage over kmeans in that it can help reveal the something of true grouping in your data.

The `hclust()` function wants a distace matrix as input. We can get this from the `dist()` function.

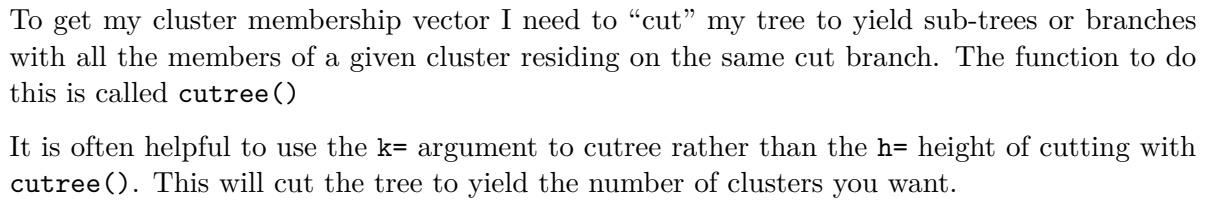
```
d = dist(x)

hc = hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=10, col="red")
```

[illegible]

Principle Component Analysis (PCA)

The base R function for PCA is called `prcomp()` Let's play with some 17D data (a very small dataset) and see how PCA can help.

PCA of UK food data

Import the data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
head(x)
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139

Q1. How many rows and columns are in your new data frame named `x`? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  5
```

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93


Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

I prefer the `row.names=1` argument because it sets a specific column as the row names while the `x[, -1]` code will delete the first row after setting them to the row names. The `row.names` function is more robust because you can run as many times and it will produce the same output, while the `x[, -1]` code will continuously delete rows as you run it repeatedly.


Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Change the `beside` argument to `false` will result in the following plot.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The code is cross comparing different countries by their food consumption. If points lie on the diagonal it means the both countries consume the same amount in that specific food.

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The main differences are their elevated consumption of fresh potatoes, lower alcohol consumption, and lower fresh fruit consumption.

```
pca = prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00


PCA plot

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14


Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
plot(pca$x[,1],pca$x[,2], col=c("orange","red","blue","darkgreen"),pch=15)
```




Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange","red","blue","darkgreen"))
```



Below we can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 67 29 4 0
```

Summarized the variances with respect to the principal component using the `summary()` function and then using `barplot()`.

```
z <- summary(pca)
z$importance
```


	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	4.188568e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



Plotting the influences of original variables against the principal components using the `barplot()` function which gives us loading scores.


```
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?


Fresh potatoes with the most positive loading score and soft drinks with the most negative loading score. PC2 tells us

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



information also can be summarized using the `biplot()` function.

```
biplot(pca)
```

New data of gene expression.

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

There are 10 samples and 100 genes.

```
ncol(rna.data)
```

```
[1] 10
```


```
nrow(rna.data)
```

```
[1] 100
```

PCA and plot the results.

```
## Again we have to take the transpose of our data  
pca <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2  
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



summary of how much variation in the original data each PC account for.

```
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642

Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251
	PC8	PC9	PC10				
Standard deviation	0.62065	0.60342	3.348e-15				
Proportion of Variance	0.00385	0.00364	0.000e+00				
Cumulative Proportion	0.99636	1.00000	1.000e+00				

barplot summary of this Proportion of Variance

```
plot(pca, main="Quick scree plot")
```




generate our own scree-plot.

```
## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```




PCA plot a bit more attractive and useful...

```
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
      xlab=paste0("PC1 (", pca.var.per[1], "%)"),
      ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```




We could use the ggplot2 package here.

```
library(ggplot2)

df <- as.data.frame(pca$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```




adding a condition specific color and sample label aesthetics.

```
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)

p
```




some spit and polish

```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data